

Report 4: groupy

Gabriele Morello

October 2022

1 Introduction

In this laboratory we implemented a simplified version of the Chord protocol, a protocol for a peer-to-peer distributed hashing table. In this protocol we will have a ring structure where each node know his successor and his predecessor, and of course his own data, we will not have a full routing table as specified by the Chord protocol. Data is stored with a key and a value, each node is responsible for a segment of keys.

2 Main problems and solutions

We have four modules: key, storage, node and test. key manages keys utilities like creation and checking if a keys belongs in an interval, storage manages the store field in the node: create a storage, lookup of a value given the key, splitting and merging of storage. The node module has 4 versions that add the functionalities described below:

The first module allows the creation of a network with no failure handling where we can only add nodes.

The second module node2 add the possibility to store data in the DHT, we introduce there the Store variable and two commands: add and lookup, respectively to add and retrieve a value in the DHT, at this stage I also creates the storage module to provide storage functions.

The third module node3 covers the first optional task, the goal is to handle failures, it does so by adding a monitor to each node for the successor and a new variable to store the successor of the successor, in case of failure of the successor a node will connect to the successor of the successor. More details will be explained later.

2.1 node1

In node1 I had to implement the stabilize procedure to update links, it works by sending a request to a possible successor and expecting a status response that states its Precedent node, based on what the response is the node can either

notify the successor to add him as predecessor, or send a new request to the predecessor of the predecessor.

The check is done by `key:between`. A stabilize call is made every second. The request function was given. `notify` is called when another node wants to be the predecessor, before adding it we still have to check that the request is valid, in case the predecessor is null or its Id is lower (in module) than the requesting node we add the new node otherwise we keep the old predecessor. `start` and `init` are given but in `connect` I had to specify the returned value.

To make this module work I had to create a key module with `key:generate()` and `key:between(Key, From, To)`, the first returns a random key, the second checks if a key is inside an interval `(From, To]` but intervals are modular because we are operating on a ring.

2.2 node2

This module extends `node1` by adding a storage in each node and the possibility to add and retrieve values by means of their key.

To achieve this I created the storage module, there we can find `add` to add a key-value pair, `lookup` to retrieve a value by means of its key, `split` to separate a storage given an interval of keys, `merge` to append two storages.

In the `node2` module we were given two new messages `add` and `lookup`. I had to implement the `add` function, this function adds a value to the storage but before it checks if the value belongs to the node with `keys:between`, in case it does we add it to the storage otherwise we say to the successor to take care of it.

`lookup` retrieves a value from the storage and it is quite similar to `add`, we check if we are in charge of that value if we are not we tell to the successor to `lookup` in its storage.

The hard part of this module was the `handover` function, what this function does is to give responsibility of keys to a new node that joined the network. The message was given I had to implement `handover` and `update` `notify` with some new features: it has to return a list of values that the new node has to take care of. We select these values with the `handover` function. `handover` uses `storage:split` to make two lists one stays in the node and it will be returned, the other will be sent with a `handover` message and will be merged with `store` in the called node.

2.3 node3

This module extends node2 and covers the first extra point: handling failures, each has to be monitored and we have to restore links to avoid interruptions, we need to keep track of the successor of the successor so in case of failure we know who will be our successor (Next variable), note that if two subsequent nodes crash our DHT will stop working, we can update Next with the status message. In stabilize we need to always return the successor and Next, we also have to take care of the case when a new node should be our successor, we need to demonitor the old node and start monitoring the new one. In the rest of the program we add monitor during the connection phase when we receive back the reference, in notify after calling handover in both the nil case and the existing predecessor case. The drop function is called also in notify in case we had an old predecessor.

When we receive the DOWN message we handle it with a down function that covers the two cases of predecessor crash and successor crash, in the first case it just leaves predecessor as nil while in the second case it starts monitoring Next, it calls stabilize and Next becomes the new successor.

3 Conclusions

This last homework was tricky not because of the Erlang code that was quite guided thanks to the skeleton but because it required a big effort to understand the protocol and how all the parts work, it was also very easy to make mistakes between variables because it was hard to keep track of the flow. Thanks to this assignment I really understood how a DHT works and I hope that this will help me for the exam.